Grep Is Dead: How I Made Claude Code Actually Remember Things

Artem Zhutov @ArtemXTech

1d • 1 tweets

Tweet Share Export

Every conversation with Claude Code starts from zero. Here's how I fixed that with a local search engine and a skill that loads your full context before you type a single word.

Every conversation with Claude Code starts from zero. I had 700 sessions in 3 weeks and I don't remember what was happening back then. I'm losing control in terms of what's happening.

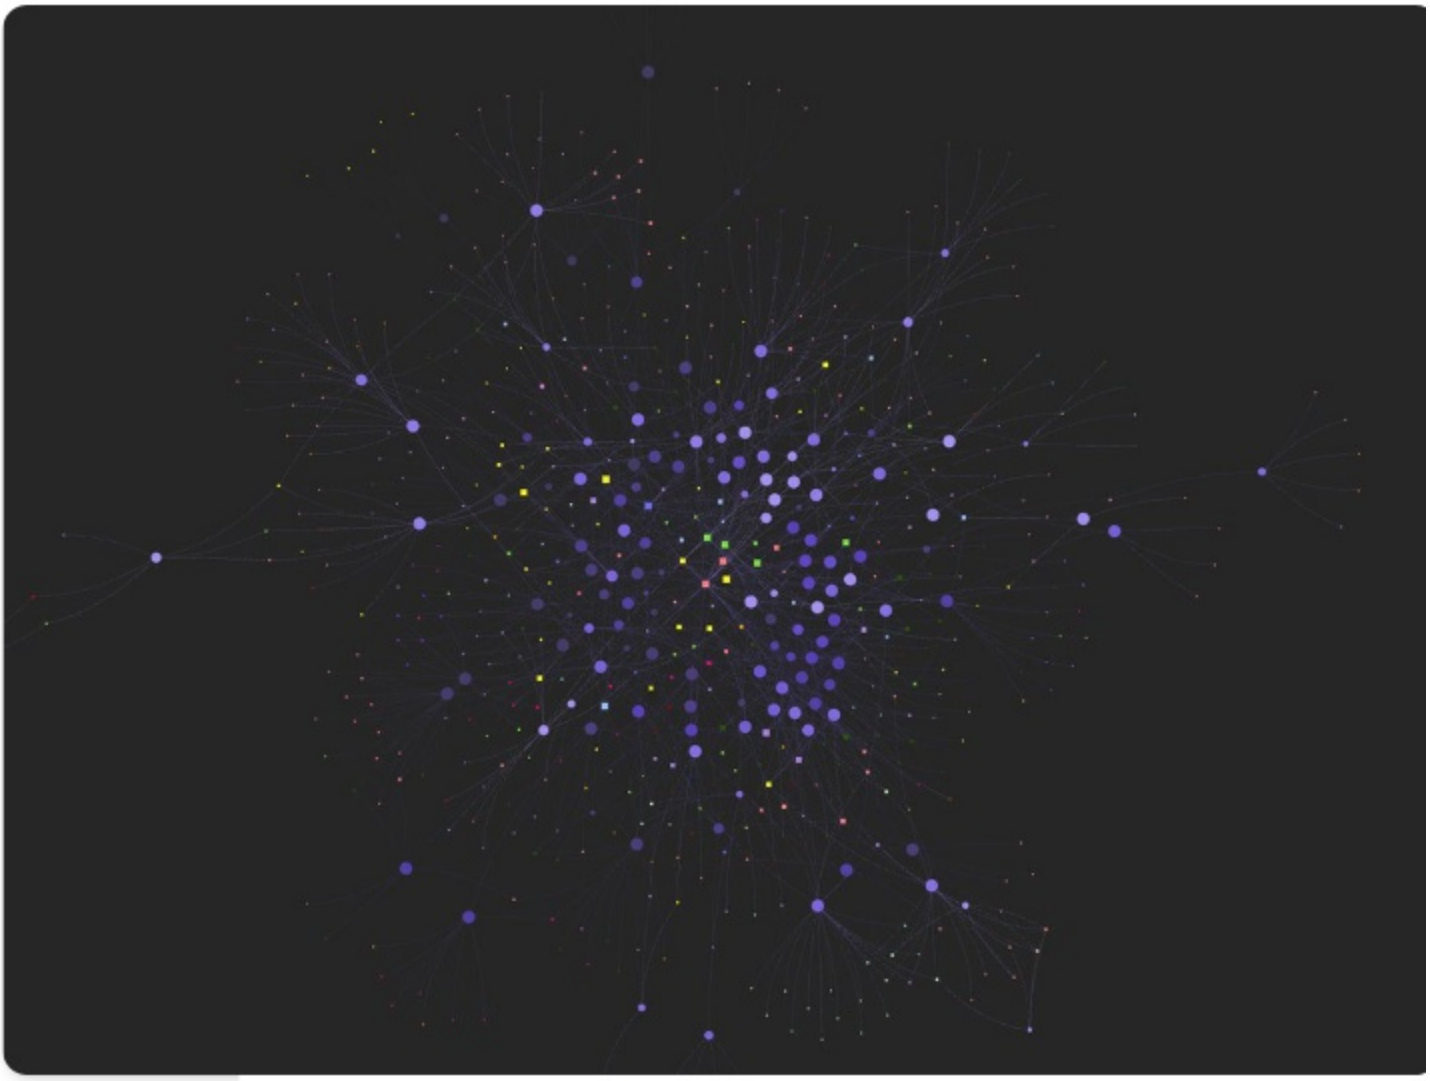

I just open a new terminal and then what's going on. I need to somehow find all this context. What was the project, what are the decisions. I need to start from zero all the time.

It gets worse mid-session. When we hit context limit at 60%, we need to compact or hand off. Then half of the decisions have been lost. Or even worse if I want to continue next day, I don't remember what was happening back then.

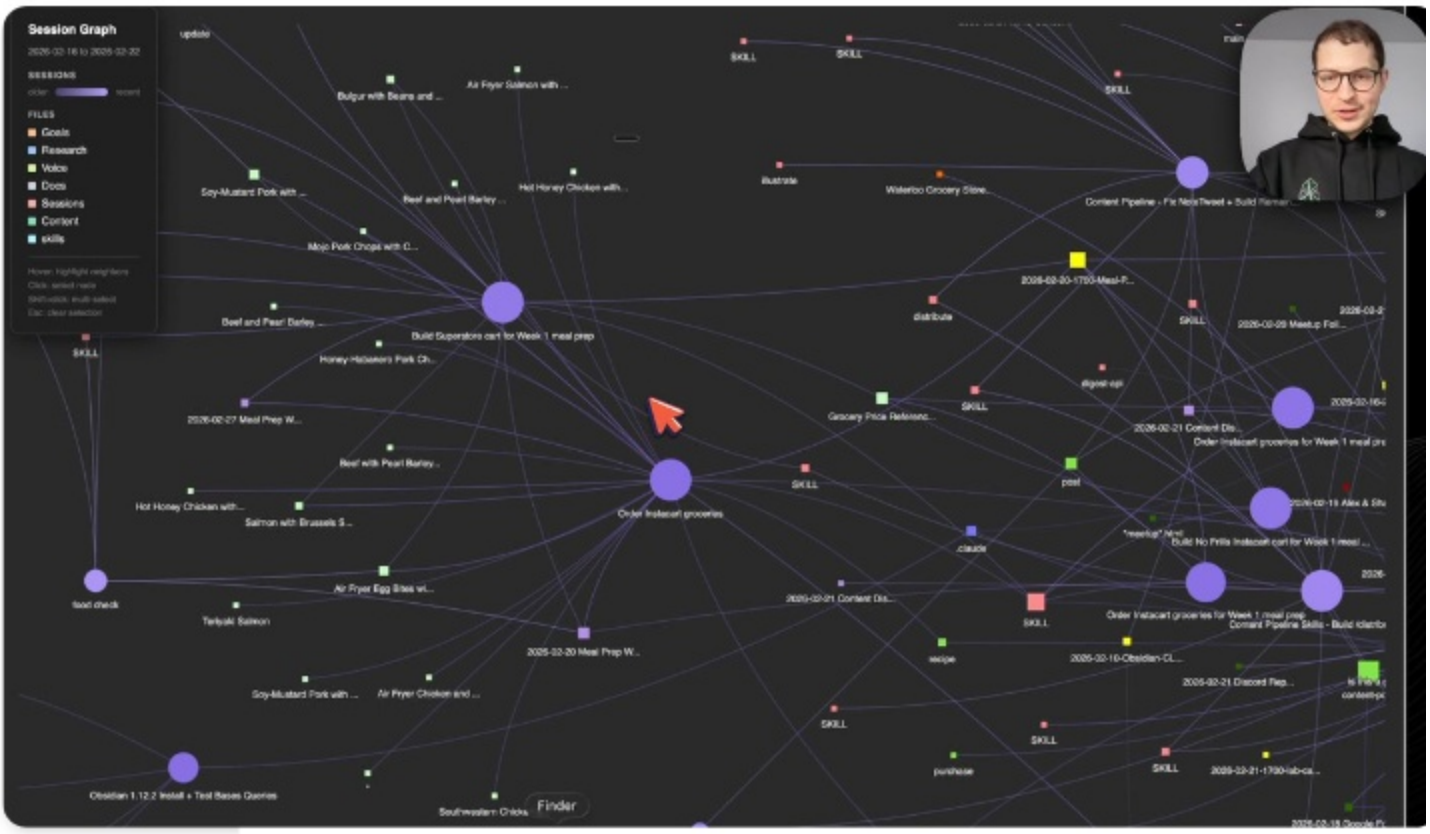

The current paradigm of grepping over files doesn't scale. So I plugged QMD into my vault, a search engine by Tobias Lutke @tobi , CEO of Shopify. Made in Canada.

The whole memory system I'm about to show you - it's a skill. Install it in 2 minutes and Claude Code already knows how to use it.

QMD: A Local Search Engine for Your Vault

QMD is a local search engine for your knowledge base. It indexes your Obsidian @obsdmd vault and finds anything in under a second.

For each vault folder I have a QMD collection. Notes, daily entries, sessions, transcripts. One-to-one mapping. For each collection I can perform a focused search.

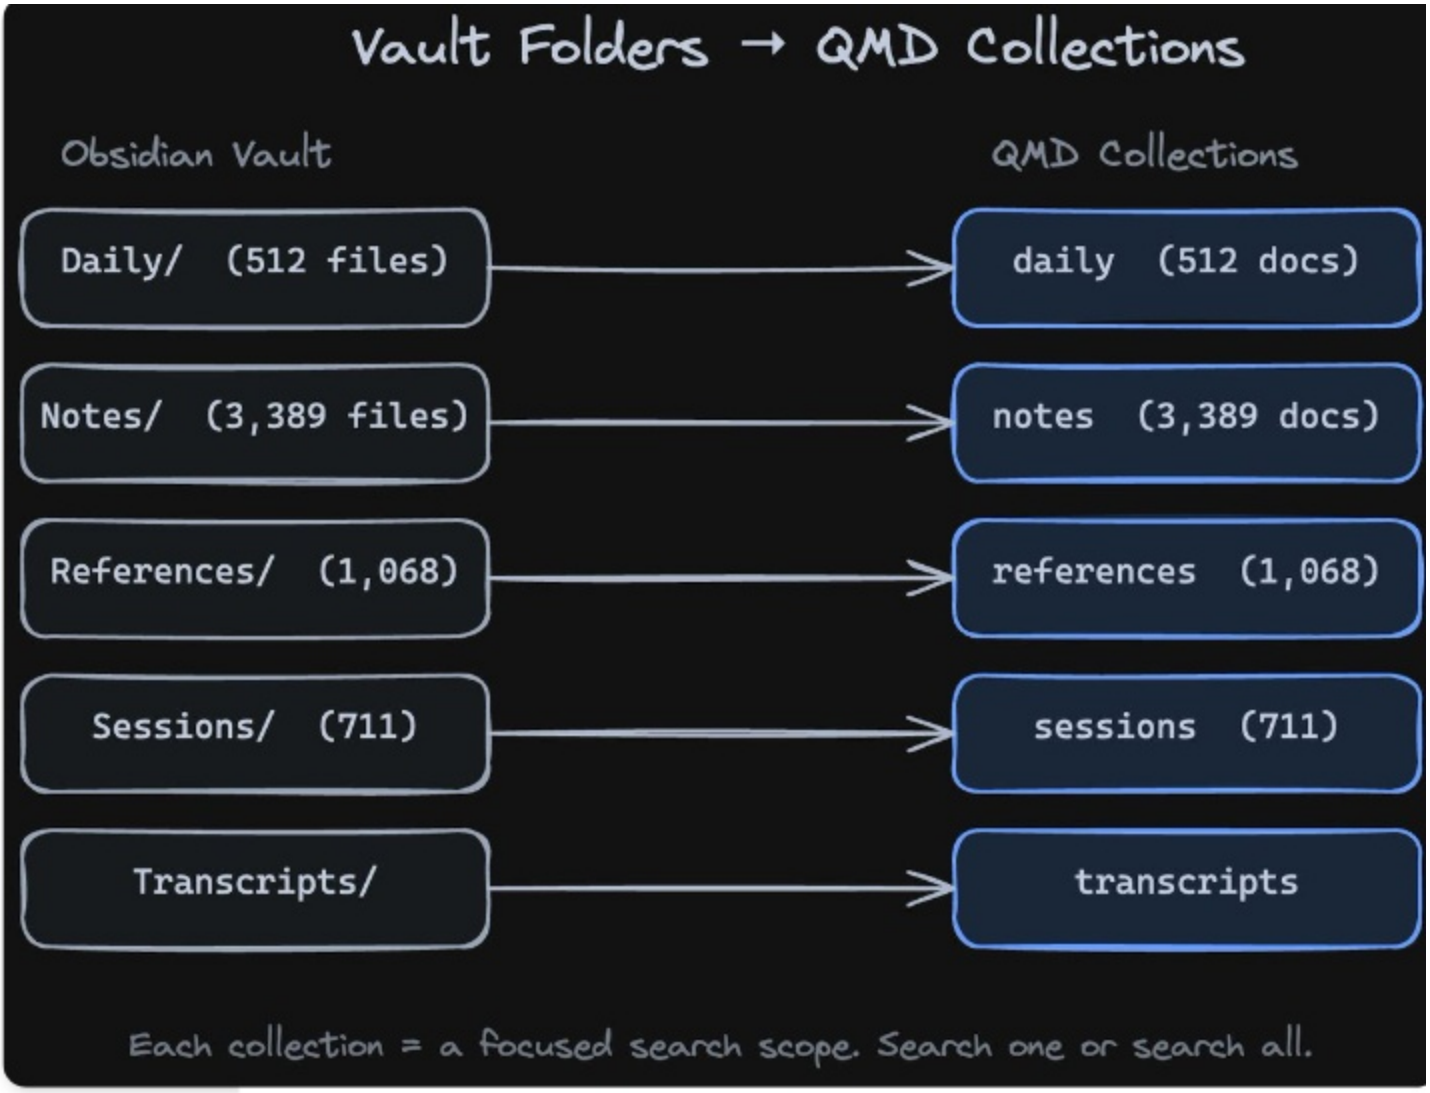

That's it. One command, your vault is searchable.

But searchable how? The default way Claude Code @claudeai searches is brute force - it sends a Haiku sub-agent to grep through every file. I tried it (watch): searched for my notes about different graph approaches the normal way. It took 3 minutes. I was bored, scrolling Twitter while waiting. And the results? 300 files, not great.

QMD search was instant. Better results. Way fewer tokens. You don't have to use sub-agents to do that. The difference: grep matches strings, QMD ranks by relevance.

Grep vs BM25 vs Semantic

QMD gives you three search modes. BM25 (qmd search) is deterministic full-text search. Like grep, it matches exact keywords. Unlike grep, it scores each file: how often does the word appear, and how rare is it across all your documents? A short note mentioning "sleep" five times scores higher than a 10,000-word file where "sleep" appears once. No AI, no embeddings - just math. Semantic (qmd vsearch) uses embeddings to find meaning - you can search for a concept even if the exact words aren't there. Hybrid (qmd query) combines both.

I searched for "sleep" across my vault. Here's the difference.

Benchmark Results

Approach	Query	Results	Time	Signal Quality
grep	"sleep"	203 files	2.8s	Wall of noise - Research docs, workshops, templates
BM25	"sleep"	5 results	0.65s	Relevant - sleep experiments, diagnostics, goals
Semantic	"couldn't sleep"	5 results	8.7s	Finds meaning - sleep patterns, schedule, energy
Hybrid	"couldn't sleep"	5 results	35s	Best ranking - "Sleep interrupted at 3am" at 90%

Key Observations

grep: 203 files, mostly noise. First results are Research docs that happen to mention "sleep" in passing (Asar Patching, Discord API, Terminal Tab Catalog). No ranking, no relevance.

BM25: Cuts to 5 ranked results. Finds the actual sleep-related files (experiments, diagnostics, issues). Fast at 0.65s. But limited to exact term matching.

Semantic (vsearch): Understands "couldn't sleep" as a concept. Finds the braindump about sleep patterns, the "Get to bed before 11:00" goal, the Evening Review Protocol - none of which necessarily contain "couldn't sleep" literally.

Hybrid (query): Best ranking - puts "Sleep interrupted at 3am" at 90% score (the most relevant doc). Also surfaces the daily note where you literally "couldn't fall back asleep." Slowest at 35s though.

Finder

Grep found 200 files. All over the place. It finds all the files which contain the word "sleep." It even finds sleep() - a programming command that pauses code execution. Nothing to do with actual sleep That's the problem with string matching.

BM25 (2 seconds): a reflection about sleep quality. Experiment with tracking sleep fragmentation patterns. Sleep interrupted at 3am. Much better.

But qmd search "insomnia" returns zero results. The word doesn't exist in the vault.

Semantic search: I type qmd vsearch "couldn't sleep, bad night" and it found a bedtime discipline goal I set years ago. It goes beyond the keywords and explores the meaning. Four of five results don't contain the search words at all. Grep could never find them.

Hybrid (qmd query "couldn't sleep, bad night"): ranks sleep quality improvement at 89%, sleep interrupted at 3am at 51%, health sleep optimization at 42%. The best ranking of all.

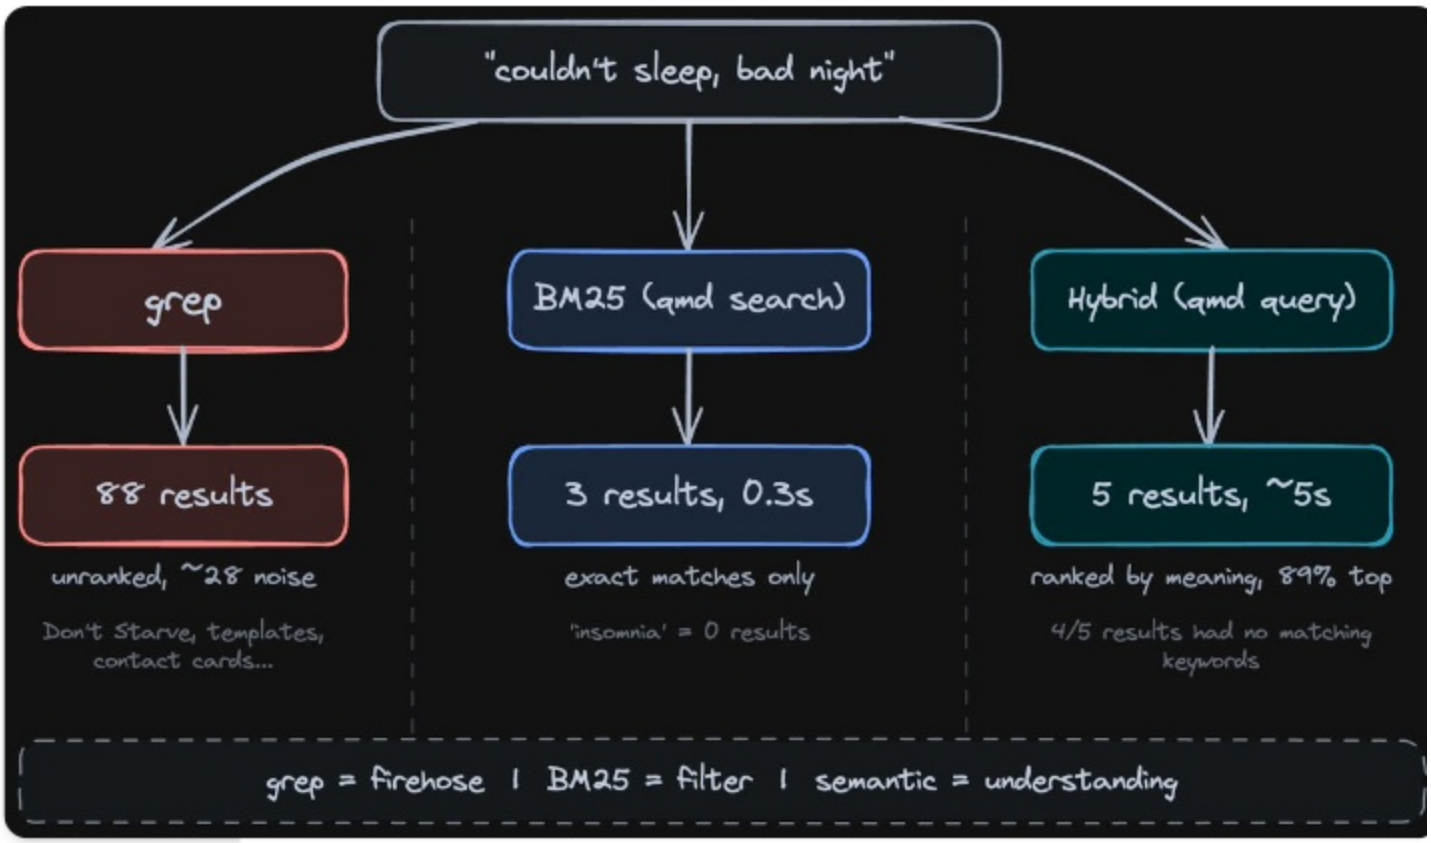

Start with BM25. Fast and handles structured notes. 80% of searches. Add semantic for transcripts and braindumps, where you'd never search for those exact words.

What semantic search actually surfaces

I searched my daily notes: "find the days when I was happy and what was the reason."

This is a non-trivial query. Claude adapted it and ran multiple searches:

Found semantically relevant connections across months of daily notes.

The pattern: my happiest days are when I ship something and I had good sleep recovery, like sauna o 9 hours of sleep.

Then it surfaced something I had completely forgotten. Back in October when I was writing my PhD thesis and I was about to give up. I needed to do it day by day, come in and write something. But what I realized back then is that I just need to see it through discomfort instead of escaping to quick fixes.

I didn't remember writing that. I didn't expect the search could surface this. But there it was, the exact citation.

/recall - load context before you start

/recall is a Claude Code skill that sits on top of QMD. It loads context before you start working - instead of explaining to Claude what you were doing, you tell it to recall.

It has three modes: temporal (scan your session history by date), topic (BM25 search across your QMD collections), and graph (interactive visualization of sessions and files).

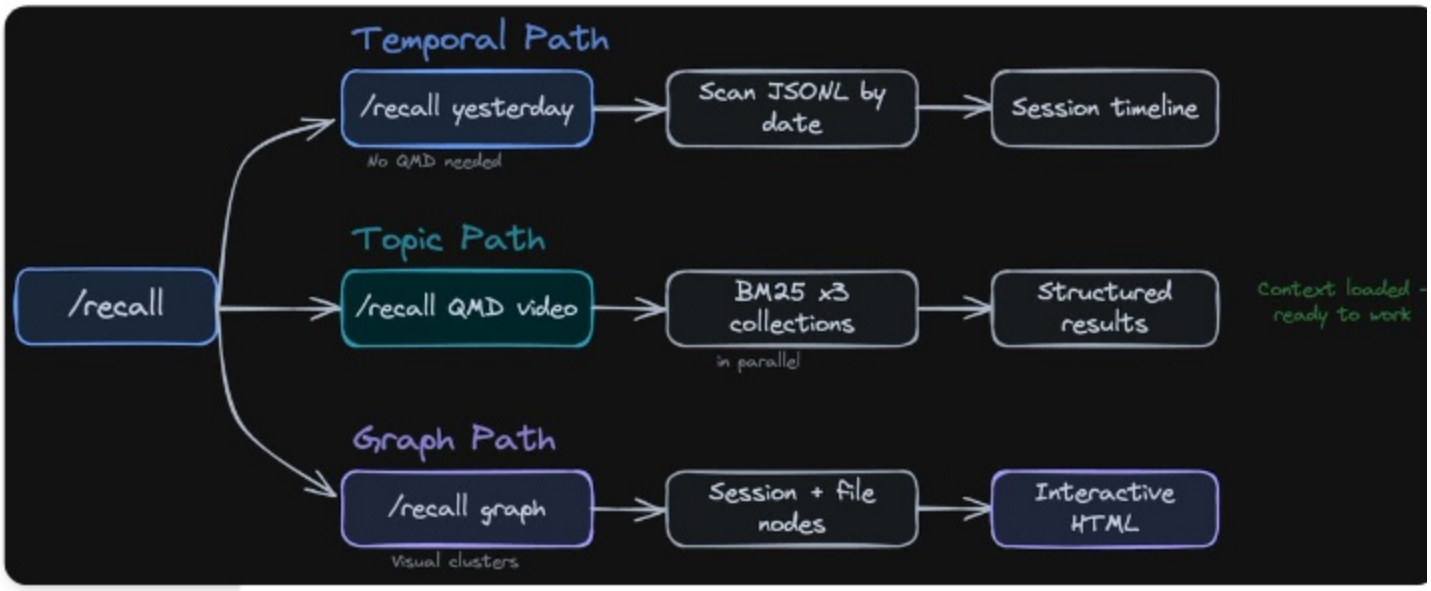

/recall yesterday

/recall topic graph

/recall graph last week

Yesterday reconstructed 39 sessions from one day. Timeline, number of messages in each session, what was done when.

/recall yesterday yes I want

okay we... ⌘1 | find he... ⌘2 | /recall... ⌘3 | i wan... ⌘4 | ~/artem-... ⌘5

● Yesterday (Feb 26) - Full Timeline

Early Morning (01:00-03:30)

Mac Mini & OpenClaw cleanup (01:12) - Removed duplicate Screen Studio apps, deleted outdated apps (postiz, pi-search, dns logger), updated OpenClaw to v2026.2.24. Found Twitter evidence of people getting banned for using Anthropic API keys with OpenClaw on Discord.

HAR Chrome Extension (01:29) - Built a Chrome extension to record HAR files on hotkey. Fixed MIME type handling (data:application/json causing .json instead of .har). Working and documented in CLAUDE.md.

Weekly Plan & Session Triage (01:49) - Triaged 34 in-progress sessions: closed 6 zombie sessions, parked 8, condensed to 6 active core. Added parked property to session template. Launched blind-spot analysis (3 agents) - identified say-do gap, stale content, B2B Workshop only 8 days away.

Albert Safin NotebookLM (03:06) - Explored Albert Safin's 7 PDFs + 17 videos on self-discipline. Fixed NotebookLM CLI Unicode encoding issues (6 places). Key moment: shared $20K from ~20 customers in 2 months - shifts narrative from "tools aren't valuable" to "distribution is the bottleneck."

Morning (05:25-08:20)

QMD Plugin Update & Reindex (05:25) - Updated QMD from v1.0.5 to v1.1.0. Benchmarked: BM25 instant, vector 2.4s, reranking 14.4s. Re-indexed vault (25 days stale) - 2,883 chunks from 794 docs in 3m 13s on Metal.

Claude Proxy for Open Notebook (05:47) - Built ~/projects/claude-proxy/ to expose OAuth as HTTP API (no ANTHROPIC_API_KEY). Registered anthropic:manual credential. Working.

GraphRAG-SDK vs GraphRAG-Lite (06:03-07:53) - Tested FalkorDB (requires Docker, team prefers none). Found GraphRAG-Lite alternative. Built graph extraction: 50 notes → 398 nodes/259 edges in 2 min. Hit concurrency issues at scale - pivoted to NetworkX (no Docker).

QMD Twitter Signal Search (05:34) - Found Toby Lutke (Shopify) created QMD. Created real vault demo: searched "food cravings" across 513 daily notes - BM25 found contextual matches in 5.3s, grep failed on synonyms.

Video Research + Content Cleanup (06:46) - Cleaned 32 stale co files (marked published/done/archived). Created Twitter X growth das

Finder

Topic searched for "QMD video" across sessions and notes. Returned the dashboard, production plan

Topic searched for "QMD video" across sessions and notes. Returned the dashboard, production plan to-do list. All related files surfaced in less than a minute. Compare that to the brute force approach: tell Claude "find all information about this project" and it sends Haiku to grep through your vault for 3 minutes, burns tokens, and comes back with worse results. With /recall topic, I reconstructed the full state of the project and asked: what is the next highest leverage action?

Graph opened an interactive visualization of my whole week. Sessions as colored blobs, older ones dimmer, recent ones highlighted in purple. Files clustered by type: goals, research, voice, docs, content, skills.

Here's an example: I was exploring lunch places. I told Claude "I want to have a great lunch" and we analyzed different places to go. I stored those as activities I might want to try. A week later, I open the graph, see that session, copy those file paths into Claude Code and continue from there. The graph makes every past conversation recoverable.

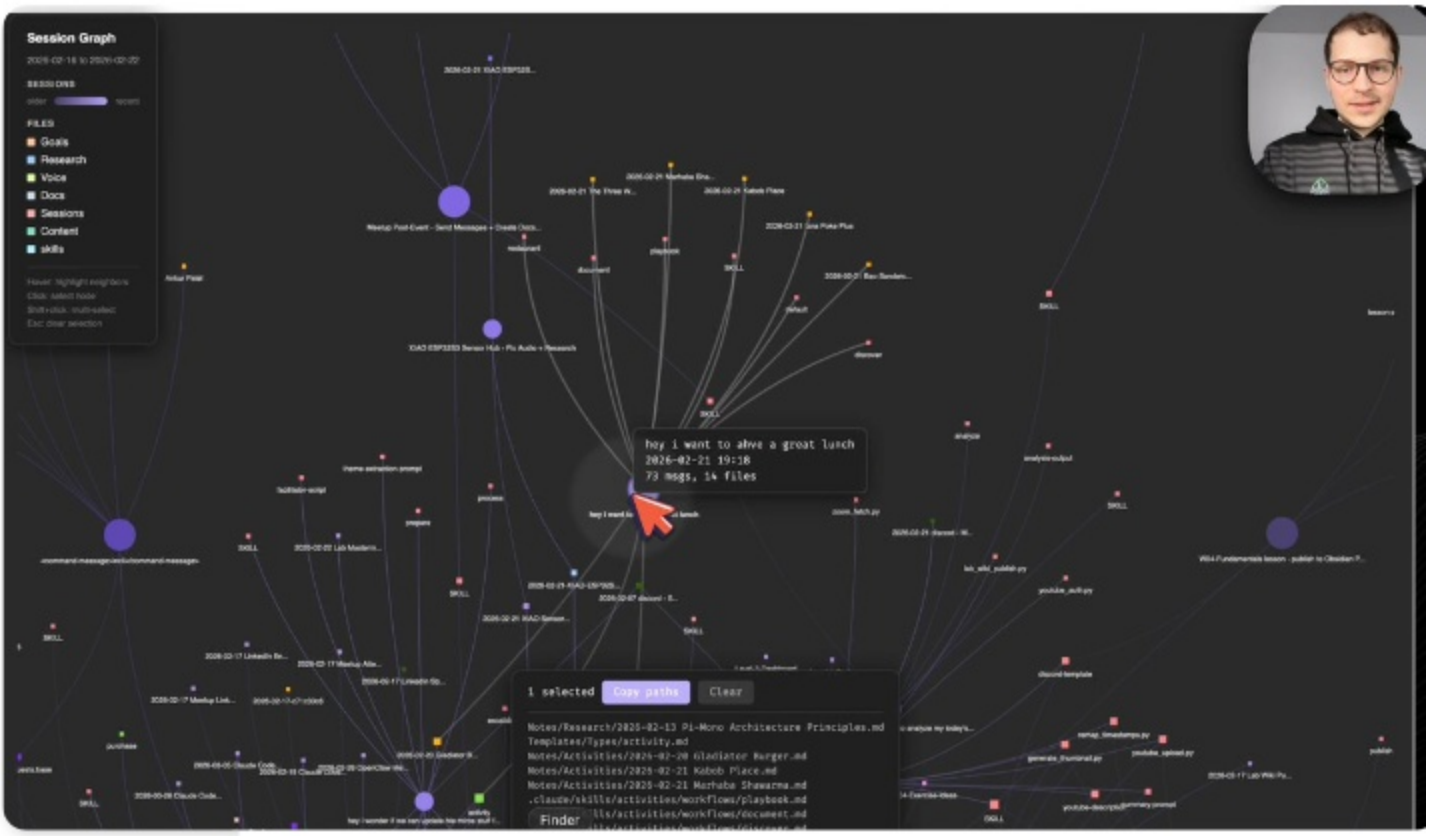

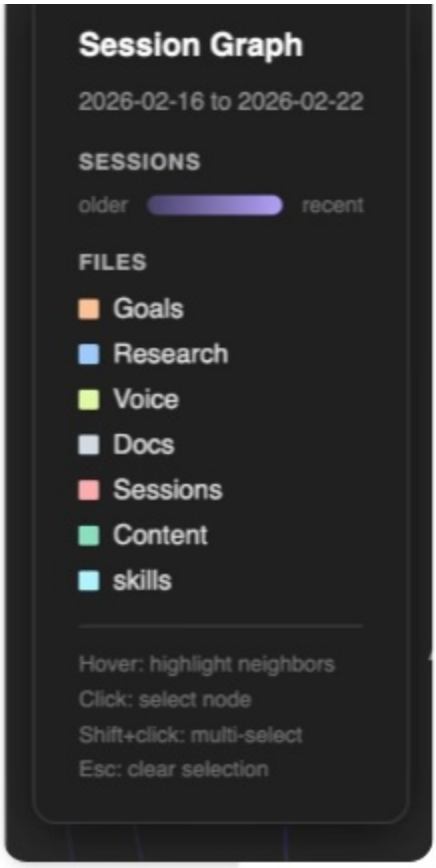

700 sessions, all searchable

Claude Code saves all your conversations as JSONL files on your computer. I had 700 sessions in the last 3 weeks. Each one has decisions, questions, context I'll need again.

The raw files have tool uses, system prompts, roles. We parse a clear markdown, the actual signal, the actual user messages, and embed this into the QMD index.

At the end of each session when I close my terminal, I have a hook which exports and embeds those sessions into QMD. So the index is always fresh. I can get back to anything I had even from today.

Finding Ideas You Never Acted On

This is the one that surprised me. I searched: "find the ideas that I have never acted on."

Having Claude synthesize the raw QMD results is very useful - it's hard to parse the actual raw results

Having Claude synthesize the raw QMD results is very useful. It's hard to parse the actual raw results yourself. Claude summarized what it found:

- October 19th - I wanted to build a PhD writing dashboard but never did it
- I had ideas for illustration-based apps but never followed through
- I had an idea to record a screen recording about my full Obsidian workflow but never committed

```
● This is a great demo of hybrid search doing something grep could never do.

  You queried "ideas I had but never acted on" - none of these files contain
  that phrase literally. But the hybrid search understood the concept and found:

  1. Oct 18 morning check-in (88%) - You wanted to build a PhD writing dashboard
   but never did it. Classic "idea I had but didn't act on."
  2. Nov 9 morning check-in (50%) - You planned job applications + YouTube
  transcripts but fell into a hole and never followed through. Pattern of
  intentions vs execution.
  3. Jul 20 morning check-in (46%) - You had an idea to record a screen
  recording showing your full Obsidian workflow - capture, voice, sync. Another
  idea that was captured but apparently not acted on.
  4. Nov 10 daily tasks (38%) - Three pending items: voice note habits, Discord
  community setup. The note literally says "been pending since Nov 4."
  5. Nov 24 morning check-in (35%) - Intentional daily choices reflection -
  meta-level thinking about what you want to do vs what you actually do.

  The search understood "ideas I had but never acted on" as a pattern of stated
  intentions without follow-through - and found exactly that across your daily
  notes. No keyword could find this. You'd have to read every daily note
  yourself to surface these patterns.

  For the video demo: this is probably the most compelling example of why
  semantic search matters on personal data. Grep finds words. This finds
  patterns in your behavior.
```

And it's all local - all of these embeddings live on your computer.

Tools change. Your context stays.

Your notes stop being passive. They stop being trapped in your Obsidian world. They actually start doing things. All of your notes is useful context about yourself that you can use to achieve goals in your life.

Tools change. A month from now there are going to be new models. So what. If you have your context you can make it work in any situation. Claude Code, Codex, Gemini CLI, anything.

The memory layer works as a skill across your whole stack. I use Obsidian Sync to keep my vault in

sync between my Mac and a Mac Mini that's always on. On the Mac Mini, OpenClaw runs 24/7. So I pick up my phone, open OpenClaw, and I have the same vault, same QMD index, same skills - from anywhere.

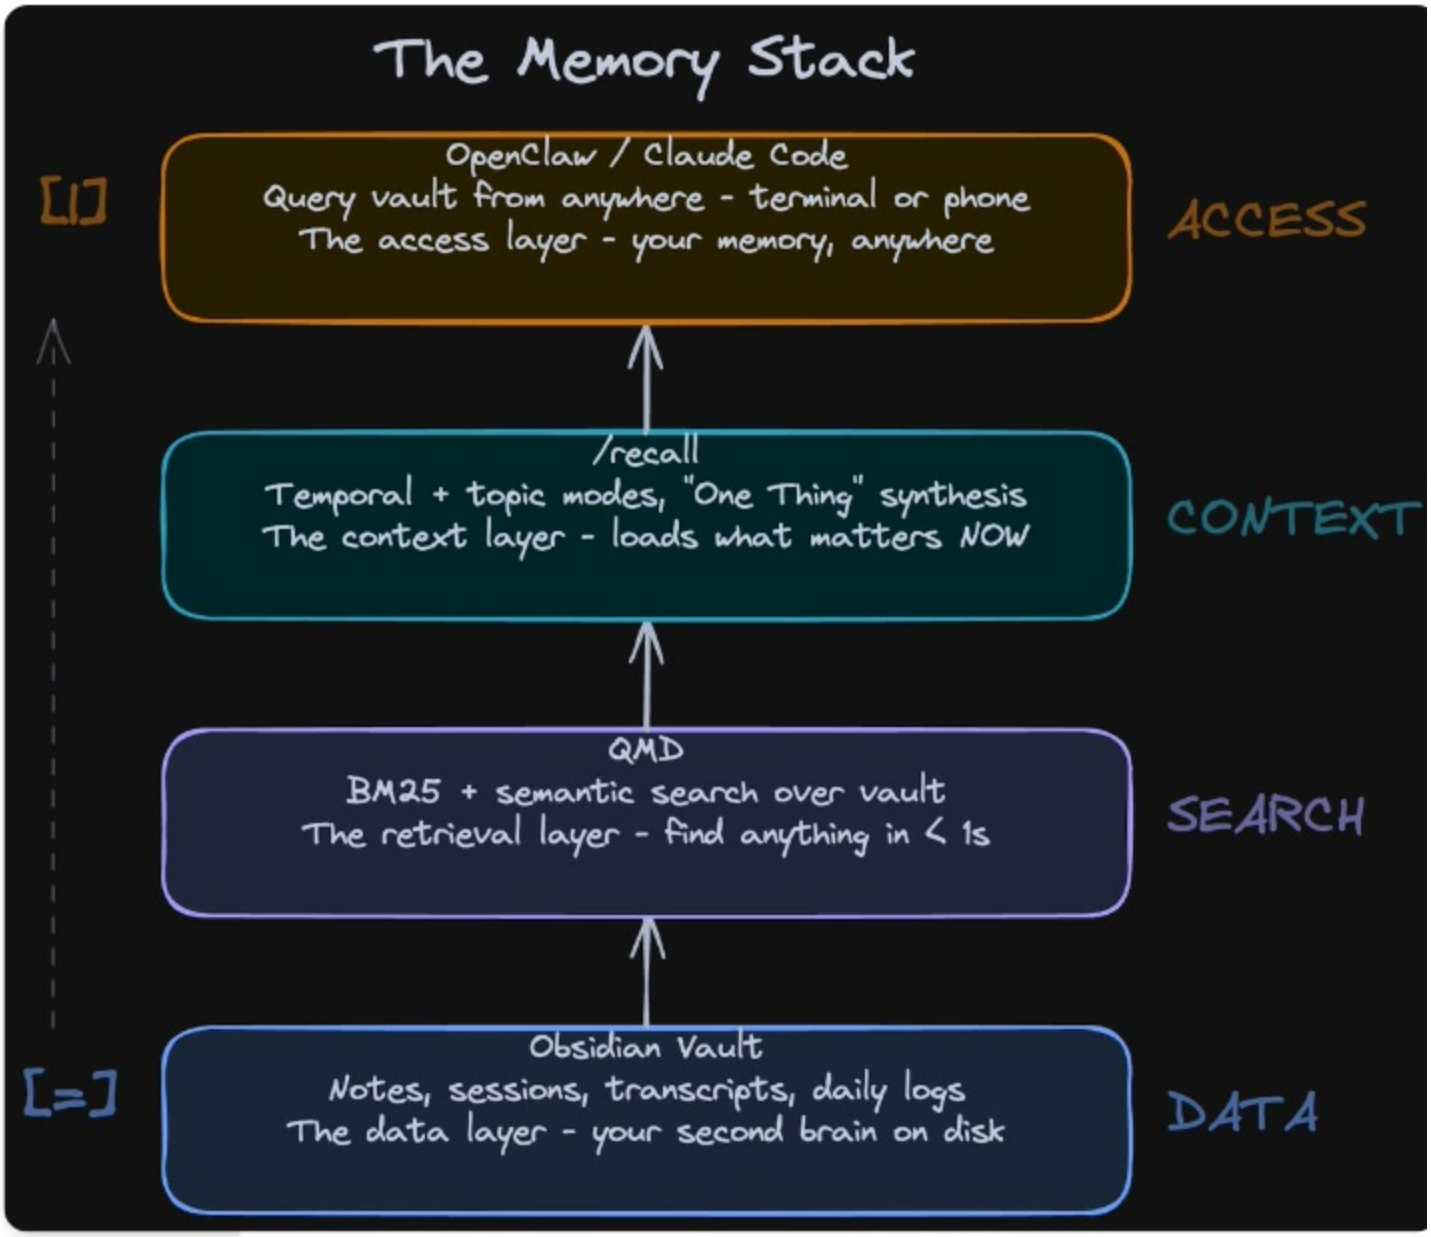

Next steps

Download the /recall skill, drop it in your .claude/skills/ folder, and you have the session pipeline and recall working today (or ask Claude to do it for you :)).

QMD - the search engine behind all of this. By Tobias Lutke.

https://github.com/tobi/qmd

Watch the full video - 42 min walkthrough with live demos.

https://youtu.be/RDoTY4_xh0s

Artem

Artem Zhutov @ArtemXTech

1d • 1 tweets

Building my Personal AI System with Claude Code + Obsidian. Physicist. Sharing what works.

Follow @ArtemXTech